

The results when the train/test split is set to the default 67/33 cut:

a. Random Forest Classifier

```
In [14]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[14]: RandomForestClassifier()
```

```
In [15]: # test the model using the test input only
prediction = model.predict(x_test)
```

```
In [16]: # calculating and printing the accuracy
model_acc = accuracy_score(y_test, prediction)*100
print(model_acc)

100.0
```

```
In [17]: # generate classification report
model_cl_rep = metrics.classification_report(y_test, prediction)
print(model_cl_rep)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	17
accuracy			1.00	27
macro avg	1.00	1.00	1.00	27
weighted avg	1.00	1.00	1.00	27

```
In [18]: # generate confusion matrix
model_cm = metrics.confusion_matrix(y_test, prediction)
print(model_cm)
```

```
[[10  0]
 [ 0 17]]
```

b. Support Vector Machines

```
In [20]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[20]: SVC()
```

```
In [21]: # test the model using the test input only
prediction_SVC = model.predict(x_test)
```

```
In [22]: # calculating and printing the accuracy
model_acc_SVC = accuracy_score(y_test, prediction_SVC)*100
print(model_acc_SVC)
```

```
77.77777777777779
```

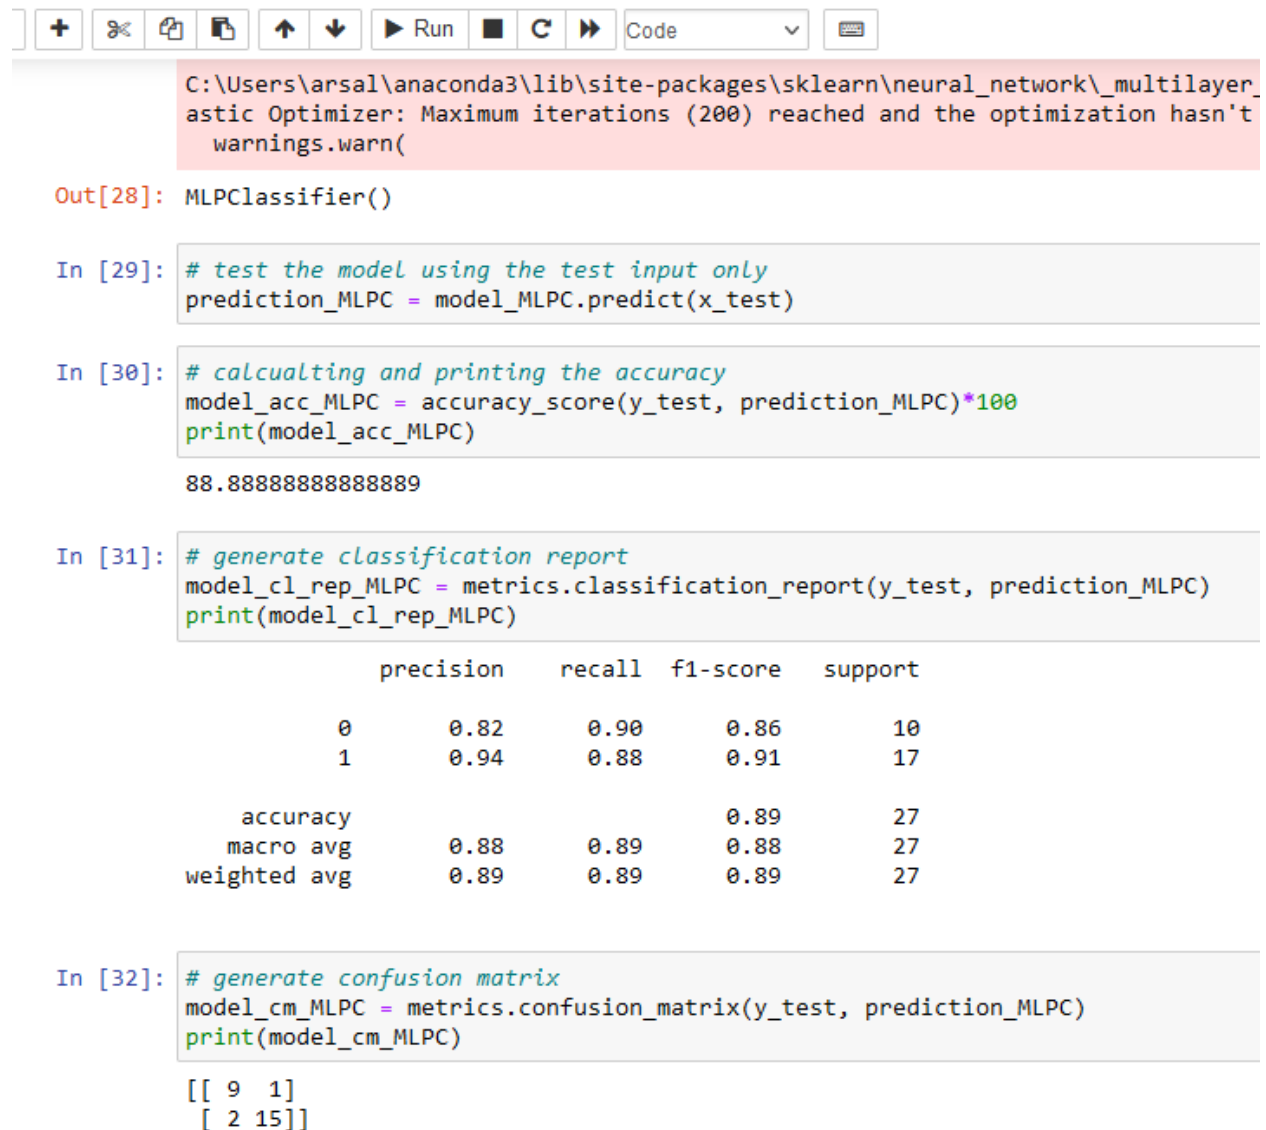
```
In [23]: # generate classification report
model_cl_rep_SVC = metrics.classification_report(y_test, prediction_SVC)
print(model_cl_rep_SVC)
```

	precision	recall	f1-score	support
0	0.70	0.70	0.70	10
1	0.82	0.82	0.82	17
accuracy			0.78	27
macro avg	0.76	0.76	0.76	27
weighted avg	0.78	0.78	0.78	27

```
In [24]: # generate confusion matrix
model_cm_SVC = metrics.confusion_matrix(y_test, prediction_SVC)
print(model_cm_SVC)
```

```
[[ 7  3]
 [ 3 14]]
```

c. Multilayer Perceptron Classifier



The image shows a Jupyter Notebook interface with a toolbar at the top containing icons for adding, deleting, and moving cells, as well as a 'Run' button. The notebook contains several code cells and their outputs.

```

C:\Users\arsal\anaconda3\lib\site-packages\sklearn\network\_multilayer_
astic Optimizer: Maximum iterations (200) reached and the optimization hasn't
warnings.warn(

Out[28]: MLPClassifier()

In [29]: # test the model using the test input only
prediction_MLPC = model_MLPC.predict(x_test)

In [30]: # calculating and printing the accuracy
model_acc_MLPC = accuracy_score(y_test, prediction_MLPC)*100
print(model_acc_MLPC)

88.88888888888889

In [31]: # generate classification report
model_cl_rep_MLPC = metrics.classification_report(y_test, prediction_MLPC)
print(model_cl_rep_MLPC)

              precision    recall  f1-score   support

      0       0.82        0.90        0.86         10
      1       0.94        0.88        0.91         17

   accuracy                   0.89         27
  macro avg       0.88        0.89        0.88         27
 weighted avg       0.89        0.89        0.89         27

In [32]: # generate confusion matrix
model_cm_MLPC = metrics.confusion_matrix(y_test, prediction_MLPC)
print(model_cm_MLPC)

[[ 9  1]
 [ 2 15]]

```

The results when the train/test split is set to 80/20 cut:

a. Random Forest Classifier

```
In [7]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[7]: RandomForestClassifier()
```

```
In [8]: # test the model using the test input only
prediction = model.predict(x_test)
```

```
In [9]: # calculating and printing the accuracy
model_acc = accuracy_score(y_test, prediction)*100
print(model_acc)
```

```
100.0
```

```
In [10]: # generate classification report
model_cl_rep = metrics.classification_report(y_test, prediction)
print(model_cl_rep)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	10
accuracy			1.00	16
macro avg	1.00	1.00	1.00	16
weighted avg	1.00	1.00	1.00	16

```
In [11]: # generate confusion matrix
model_cm = metrics.confusion_matrix(y_test, prediction)
print(model_cm)
```

```
[[ 6  0]
 [ 0 10]]
```

b. Support Vector Machines

```
In [13]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[13]: SVC()
```

```
In [14]: # test the model using the test input only
prediction_SVC = model.predict(x_test)
```

```
In [15]: # calculating and printing the accuracy
model_acc_SVC = accuracy_score(y_test, prediction_SVC)*100
print(model_acc_SVC)
```

```
81.25
```

```
In [16]: # generate classification report
model_cl_rep_SVC = metrics.classification_report(y_test, prediction_SVC)
print(model_cl_rep_SVC)
```

	precision	recall	f1-score	support
0	0.80	0.67	0.73	6
1	0.82	0.90	0.86	10
accuracy			0.81	16
macro avg	0.81	0.78	0.79	16
weighted avg	0.81	0.81	0.81	16

```
In [17]: # generate confusion matrix
model_cm_SVC = metrics.confusion_matrix(y_test, prediction_SVC)
print(model_cm_SVC)
```

```
[[4 2]
 [1 9]]
```

c. Multilayer Perceptron Classifier

```
astic Optimizer: Maximum iterations (200) reached and the optimization hasn't
warnings.warn(
```

```
Out[19]: MLPClassifier()
```

```
In [20]: # test the model using the test input only
prediction_MLPC = model_MLPC.predict(x_test)
```

```
In [21]: # calculating and printing the accuracy
model_acc_MLPC = accuracy_score(y_test, prediction_MLPC)*100
print(model_acc_MLPC)
```

```
100.0
```

```
In [22]: # generate classification report
model_cl_rep_MLPC = metrics.classification_report(y_test, prediction_MLPC)
print(model_cl_rep_MLPC)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	10
accuracy			1.00	16
macro avg	1.00	1.00	1.00	16
weighted avg	1.00	1.00	1.00	16

```
In [23]: # generate confusion matrix
model_cm_MLPC = metrics.confusion_matrix(y_test, prediction_MLPC)
print(model_cm_MLPC)
```

```
[[ 6  0]
 [ 0 10]]
```

Q 2.4 Results after removing height and hair length

a. Random Forest Classifier

```
model.fit(X_train,Y_train)
```

```
Out[27]: RandomForestClassifier()
```

```
In [28]: # test the model using the test input only
prediction = model.predict(x_test)
```

```
In [29]: # calculating and printing the accuracy
model_acc = accuracy_score(y_test, prediction)*100
print(model_acc)
```

```
100.0
```

```
In [30]: # generate classification report
model_cl_rep = metrics.classification_report(y_test, prediction)
print(model_cl_rep)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	6
1	1.00	1.00	1.00	10
accuracy			1.00	16
macro avg	1.00	1.00	1.00	16
weighted avg	1.00	1.00	1.00	16

```
In [31]: # generate confusion matrix
model_cm = metrics.confusion_matrix(y_test, prediction)
print(model_cm)
```

```
[[ 6  0]
 [ 0 10]]
```

b. Support Vector Machines

```
In [33]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[33]: SVC()
```

```
In [34]: # test the model using the test input only
prediction_SVC = model.predict(x_test)
```

```
In [35]: # calculating and printing the accuracy
model_acc_SVC = accuracy_score(y_test, prediction_SVC)*100
print(model_acc_SVC)
```

```
81.25
```

```
In [36]: # generate classification report
model_cl_rep_SVC = metrics.classification_report(y_test, prediction_SVC)
print(model_cl_rep_SVC)
```

	precision	recall	f1-score	support
0	0.80	0.67	0.73	6
1	0.82	0.90	0.86	10
accuracy			0.81	16
macro avg	0.81	0.78	0.79	16
weighted avg	0.81	0.81	0.81	16

```
In [37]: # generate confusion matrix
model_cm_SVC = metrics.confusion_matrix(y_test, prediction_SVC)
print(model_cm_SVC)
```

```
[[4 2]
 [1 9]]
```


c. Multilayer Perceptron Classifier

```
In [38]: # changing the classifier as required
model_MLPC = MLPClassifier()
```

```
In [39]: # training the model using the training data
model_MLPC.fit(X_train,Y_train)
```

```
Out[39]: MLPClassifier()
```

```
In [40]: # test the model using the test input only
prediction_MLPC = model_MLPC.predict(x_test)
```

```
In [41]: # calculating and printing the accuracy
model_acc_MLPC = accuracy_score(y_test, prediction_MLPC)*100
print(model_acc_MLPC)
```

```
62.5
```

```
In [42]: # generate classification report
model_cl_rep_MLPC = metrics.classification_report(y_test, prediction_MLPC)
print(model_cl_rep_MLPC)
```

```
...
```

```
In [43]: # generate confusion matrix
model_cm_MLPC = metrics.confusion_matrix(y_test, prediction_MLPC)
print(model_cm_MLPC)
```

```
[[ 0  6]
 [ 0 10]]
```

QUESTION 3

a. Monte Carlo Cross Validation

```
In [4]: # selecting a classifier
# model = RandomForestClassifier()

# create a Gaussian classifier
# model = GaussianNB()

# create a Decision tree classifier
model = DecisionTreeClassifier()

In [5]: # separating independent and dependent variables from the data frame
x = list(zip(df_gender.height, df_gender.weight, beard_encd, hair_length_encd, df_gender.shoe_size, scarf_encd, eye_color_encd,
y = gender_encd

In [9]: # for shuffle split model
mc = ShuffleSplit(n_splits = 5, test_size = 0.33, random_state = 7)

In [10]: print("Cross validation scores with Monte Carlo Cross Validation")
cross_val_score(model, x, y, cv = mc).mean()

Cross validation scores with Monte Carlo Cross Validation

Out[10]: 0.9517241379310345
```

b. Leave P-Out Cross Validation

```
In [3]: # converting the string based inputs into integer as the model only understands integers
labels = preprocessing.LabelEncoder()
beard_encd = labels.fit_transform(df_gender.beard)
hair_length_encd = labels.fit_transform(df_gender.hair_length)
scarf_encd = labels.fit_transform(df_gender.scarf)
eye_color_encd = labels.fit_transform(df_gender.eye_color)
gender_encd = labels.fit_transform(df_gender.gender)

In [4]: # selecting a classifier
# model = RandomForestClassifier()

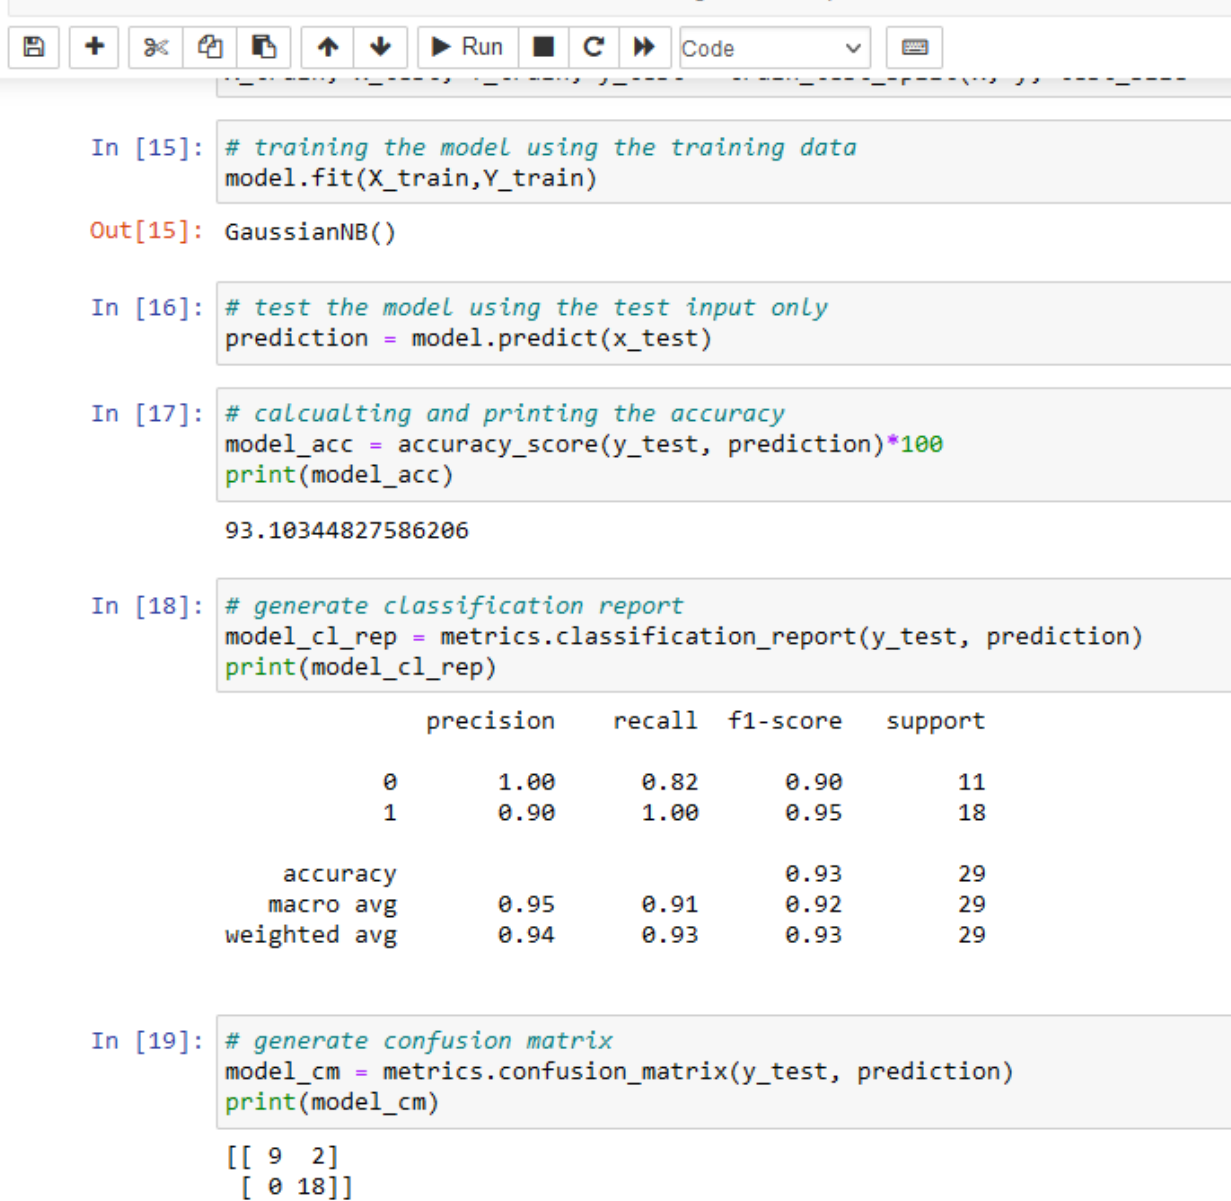
# create a Gaussian classifier
# model = GaussianNB()

# create a Decision tree classifier
model = DecisionTreeClassifier()

In [5]: # separating independent and dependent variables from the data frame
x = list(zip(df_gender.height, df_gender.weight, beard_encd, hair_length_encd, df_gender.shoe_size, scarf_encd, eye_color_encd,
y = gender_encd

In [6]: # Leave POut stuff
lpo = LeavePOut(p=2)
lpo.get_n_splits(x)
tree = RandomForestClassifier(n_estimators = 10, max_depth = 5, n_jobs=-1)
score = cross_val_score(tree, x, y, cv = lpo)
print("Cross Validation Scores are {}".format(score))
print("Average Cross Validation score :{}".format(score.mean()))

Cross Validation Scores are [1. 1. 1. ... 1. 1. 1.]
Average Cross Validation score :0.9683473389355742
```

QUESTION 4


The image shows a Jupyter Notebook interface with a toolbar at the top containing icons for saving, adding, deleting, copying, pasting, and running code. The notebook contains four input cells (In [15] to In [19]) and their corresponding outputs.

```
In [15]: # training the model using the training data
model.fit(X_train,Y_train)
```

```
Out[15]: GaussianNB()
```

```
In [16]: # test the model using the test input only
prediction = model.predict(x_test)
```

```
In [17]: # calculating and printing the accuracy
model_acc = accuracy_score(y_test, prediction)*100
print(model_acc)
```

```
93.10344827586206
```

```
In [18]: # generate classification report
model_cl_rep = metrics.classification_report(y_test, prediction)
print(model_cl_rep)
```

	precision	recall	f1-score	support
0	1.00	0.82	0.90	11
1	0.90	1.00	0.95	18
accuracy			0.93	29
macro avg	0.95	0.91	0.92	29
weighted avg	0.94	0.93	0.93	29

```
In [19]: # generate confusion matrix
model_cm = metrics.confusion_matrix(y_test, prediction)
print(model_cm)
```

```
[[ 9  2]
 [ 0 18]]
```

[End of File]